


- ② Consider a three-layer neural network for classification with output units employing softmax, trained with 0-1 signals
- (a) Derive the learning rule if the criterion function is SSE, that is,

$$J(w) = \frac{1}{2} \sum_{k=1}^K (t_k - z_k)^2$$

Derivative in terms of output unit weights:

$$\begin{aligned} \frac{\delta J}{\delta w_{kj}} &= \frac{\delta J}{\delta z_k} \frac{\delta z_k}{\delta \text{net}_k} \frac{\delta \text{net}_k}{\delta w_{kj}} \\ &= (t_k - z_k)(-1) \cdot \frac{e^{\text{net}_k}}{\sum e^{\text{net}_k}} \left(1 - \frac{e^{\text{net}_k}}{\sum e^{\text{net}_k}}\right) \cdot y_j \end{aligned}$$



Update rule for w_{kj} :

$$w_{kj} \leftarrow w_{kj} - \alpha \frac{\delta J}{\delta w_{kj}}$$

Derivative in terms of ~~output~~ hidden-to-output (y_j) $f(\text{net})$

$$\begin{aligned} \frac{\delta J}{\delta y_j} &= \sum_k \frac{\delta J}{\delta z_k} \frac{\delta z_k}{\delta \text{net}_k} \frac{\delta \text{net}_k}{\delta y_j} \\ &= \sum_k (t_k - z_k) \cdot \frac{e^{\text{net}_k}}{\sum e^{\text{net}_k}} \left(1 - \frac{e^{\text{net}_k}}{\sum e^{\text{net}_k}}\right) \cdot w_{kj} \end{aligned}$$

Derivative of loss function in terms of hidden unit weights

$$\frac{\delta J}{\delta w_{ji}} = \frac{\delta J}{\delta y_j} \cdot \frac{\delta y_j}{\delta \text{net}_j} \cdot \frac{\delta \text{net}_j}{\delta w_{ji}}$$

$$= \left[\sum_k (z_k - t_k) \cdot \frac{e^{\text{net}_k}}{\sum e^{\text{net}_k}} \left(1 - \frac{e^{\text{net}_k}}{\sum e^{\text{net}_k}}\right) \cdot w_{kj} \right] \cdot f'(\text{net}_j) \cdot x_i$$

Update rule for w_{ji} :

$$w_{ji} \leftarrow w_{ji} - \alpha \frac{\delta J}{\delta w_{ji}}$$

- (b) Repeat for cross-entropy criterion function

$$J_{ce}(w) = \sum_{k=1}^K t_k \ln \frac{t_k}{z_k}$$

$$\frac{\delta J_{ce}}{\delta z_k} = \frac{\delta}{\delta z_k} t_k \ln \frac{t_k}{z_k} = t_k \cdot \frac{1}{z_k} \cdot (-1) \frac{t_k}{z_k} = -\frac{t_k}{z_k}$$

The rest follows directly from (a)

SEE Q1
@BACK!

ANDREW DUDLEY
1001149051

CSE589 HW4

Q3

Suppose there is a one-dimensional mixture density consisting of two Gaussian components, each centered at the origin.

$$p(x|O) = P(w_1) \frac{1}{\sqrt{2\pi}\sigma_1} e^{-x^2/(2\sigma_1^2)} + (1-P(w_1)) \frac{1}{\sqrt{2\pi}\sigma_2} e^{-x^2/(2\sigma_2^2)}$$

and $O = (P(w_1), \sigma_1, \sigma_2)^T$

(a) Show that under these conditions, the density is completely unidentifiable.

If the variance of the Gaussian ~~components~~ components is equal, we get $\frac{1}{\sqrt{2\pi}\sigma} e^{-x^2/(2\sigma^2)} (P(w_1) + (1-P(w_1)))$.

Now for any value of $P(w_1)$, the density is the same, making it unidentifiable.

(b) ~~Suppose the value of $P(w_1)$ is fixed and known. Is the model identifiable?~~
 Suppose the value of $P(w_1)$ is fixed and known. Is the model identifiable?

If the prior = 0, we cannot recover σ_1 ,

" " = 1, we cannot recover σ_2

" " = 0.5, we can't distinguish between σ_1 and σ_2

otherwise, the model is identifiable.

(c) Suppose σ_1 & σ_2 are ~~known~~ known, but $P(w_1)$ is not. Is the resulting model identifiable?

Following from (a), if $\sigma_1 = \sigma_2$, then $P(w_1)$ is not identifiable.

If $\sigma_1 \neq \sigma_2$, then the model is identifiable.

Q4. On the k-means algorithm

- (1) Will the algo still work if μ_i for $i=1, \dots, C$ to the same initial value?

Yes, though in practice its results will probably depend on how empty sets are handled.

When the centroids are initialized + then a single iteration occurs, the centroids will likely change, at which point there is no longer a concern of them having been initialized at the same point.

- (2) What is the complexity of the algorithm in terms of N (the # of samples), C (the # of clusters), and T (the total # of iterations until convergence)?

$O(NCT)$

- (3) Consider the case of N data points drawn from a mixture density model w/ C normal densities with means at μ_i , $i=1, \dots, C$, respectively. After running the k-means algorithm on this data set will you get the respective mean vectors μ_i as the outcome?

You will get some vectors back, but with finite data points there is no guarantee of getting the true parameters θ , and even with infinite data points, the ground-truth parameters of the mixture density would still play a role in determining if θ is identifiable.

```

#Q1
# (A)
# # Generate training data and labels
# training_w1_data = multivariate_normal(mean_w1, covar_w1, size=(50))
# training_w1_labels = np.zeros(training_w1_data.shape[0])
# training_w2_data = multivariate_normal(mean_w2, covar_w2, size=(50))
# training_w2_labels = np.ones(training_w2_data.shape[0])
#
# # Generate test data and labels
# test_w1_data = multivariate_normal(mean_w1, covar_w1, size=(20))
# test_w1_labels = np.zeros(test_w1_data.shape[0])
# test_w1_labels = test_w1_labels.reshape((-1, 1))
# test_w2_data = multivariate_normal(mean_w2, covar_w2, size=(20))
# test_w2_labels = np.ones(test_w2_data.shape[0])
# test_w2_labels = test_w2_labels.reshape((-1, 1))
# training_data = np.r_[training_w1_data, training_w2_data]
# training_labels = np.r_[training_w1_labels, training_w2_labels]
# training = unison_shuffled_copies(training_data, training_labels)
# training[1] = training[1].reshape((-1, 1))
# test_data = np.r_[test_w1_data, test_w2_data]
# test_labels = np.r_[test_w1_labels, test_w2_labels]
# test = unison_shuffled_copies(test_data, test_labels)
# test[1] = test[1].reshape((-1, 1))

# (B) Best validation error achieved: 22%
# n*_H = 4
# (C) 10%
# (D) In part B, we fit the data too strongly to the training data. In doing so,
# the validation accuracy decreased.
# In Part C, we stopped training at the minimum value, causing us to avoid
# overfitting.

import numpy as np, math
import argparse
import pickle
import matplotlib

matplotlib.use('agg')
from matplotlib import pyplot as plt

import os, sys
from tensorflow.examples.tutorials.mnist import input_data

alpha = 0.1

class NeuralNetwork:
    def __init__(self, *args, dropout=False):
        # Layers of the neural network (doesn't include an input layer)
        self.loss = []
        self.layers = []

        self.image_count = 0

        current_wire = None
        for i, (layerType, units) in enumerate(*args):
            # For each layer defined, create the layer and attach it to the output wires
            # of the preceding layer

            if i == 0:
                # Create wires going from inputs to first hidden layer
                current_wire = Wire(units)
            else:
                new_layer = {
                    'sigmoid': sigmoidLayer(current_wire, units, dropout=dropout),
                    'softmax': softmaxLayer(current_wire, units, dropout=dropout),
                    'relu': reluLayer(current_wire, units, dropout=dropout)
                }.get(layerType, None)
                current_wire = new_layer.outputWire
                self.layers.append(new_layer)

    def train_and_validate(self, train_data, validation_data, epochs):
        global alpha
        alpha = 0.1
        for epoch in range(epochs):
            print("Epoch:", epoch)
            self.train(train_data)
            self.test(validation_data)
            alpha *= 1

    def train(self, data):
        train_loss = []
        for i, (input, label) in enumerate(zip(data[0], data[1])):
            output = self.forward_pass(input, training=True)
            errorGradient = output - label # TODO: Modify this when gradient calculation is resolved
            errorGradient = errorGradient.reshape((-1, 1))
            print("training", "output:", output, "label:", label, "errorGradient:", errorGradient)
            train_loss.append(-label.T.dot(output))
            self.backward_pass(errorGradient)

```



```

self.loss.append(sum(train_loss) / float(len(train_loss)))
plt.clf()
plt.plot(self.loss)
# plt.savefig('images/training_' + str(self.image_count) + '.png')
self.image_count += 1

def test(self, data):
    correct_count = 0
    for i, (input, label) in enumerate(zip(data[0], data[1])):
        output = self.forward_pass(input)
        print(input, output, label)
        if np.round(output) == label:
            correct_count += 1

    print(correct_count / data[0].shape[0])
    print()

def forward_pass(self, input, training=False):
    current_output = input
    for i, layer in enumerate(self.layers):
        current_output = layer.forward(current_output, training=training)
    return current_output

def calculate_error(self, prediction, label):
    return 0.5 * (prediction - label) ** 2 # TODO: Modify when gradient issue resolved

def backward_pass(self, errorGradient):
    for i, layer in enumerate(self.layers[::-1]):
        if i == 0:
            layer.outputWire.gradients = errorGradient
            layer.backward() # TODO: Modify when gradient issue resolved
        else:
            layer.backward()

def saveModel(self):
    pickle.dump(self, open("model", "wb"))

@staticmethod
def load_model():
    return pickle.load(open("model", "rb"))

class Wire:
    def __init__(self, input_dimensions):
        self.input_dimensions = input_dimensions

    def initialize(self, units, fn):
        self.dropout_proba = 0.5
        self.units = units
        self.weights = fn(self.input_dimensions + 1, units)
        self.gradients = None
        self.velocity = np.zeros(self.weights.shape)

    def generate_dropout_variables(self):
        self.dropout_vector = np.random.choice([0, 1], size=(self.units,),
                                                p=[self.dropout_proba, 1 - self.dropout_proba])
        self.dropout_matrix = np.diag(self.dropout_vector)

class Layer:
    def __init__(self, inputWire, units, dropout=False, weight_fn=lambda x, y: np.random.randn(x, y) / np.sqrt(x)):
        if not inputWire:
            raise TypeError("Must initialize layer with Wire")
        self.dropout = dropout
        self.inputWire = inputWire
        self.inputWire.initialize(units, fn=weight_fn)

        self.outputWire = Wire(units)

    def forward(self, x, training=False):
        raise NotImplementedError("Must implement forward-pass function -- forward( self, )")

    def backward(self):
        raise NotImplementedError("Must implement backward-pass function -- backward( self, )")

class sigmoidLayer(Layer):
    def __init__(self, inputWire, units, dropout=False):
        super().__init__(inputWire, units, dropout)

    def forward(self, x, training=False):
        self.inputs = x
        # Add input value of "1" to the input array for the bias weight (w_0*x_0 + ...)
        self.inputs = np.insert(x, 0, 1.0, 0)
        self.linear_outputs = self.inputWire.weights.T.dot(self.inputs)

        self.outputs = self.sigmoid(self.linear_outputs)
        if training and self.dropout:
            self.inputWire.generate_dropout_variables()
            # self.outputs = self.outputs.dot(self.inputWire.dropout_matrix)
            self.outputs = self.outputs * self.inputWire.dropout_vector / self.inputWire.dropout_proba
        return self.outputs

```

```

def sigmoid(self, x):
    # Calculate sigmoid on vector cells
    self.sig_values = 1 / (1 + np.exp(-x))
    return self.sig_values

def backward(self):
    d_p_wrt_f = self.sig_values * (1 - self.sig_values)
    if self.dropout:
        d_p_wrt_f = d_p_wrt_f * self.inputWire.dropout_vector
    self.d_E_wrt_f = np.sum(self.outputWire.gradients) * d_p_wrt_f if self.outputWire.gradients.shape[
        0] > 1 else self.outputWire.gradients * d_p_wrt_f

    d_E_wrt_x = self.inputWire.weights.dot(
        np.diag(self.d_E_wrt_f)) # Multiply outgoing gradients by the dE/df of their respective node in this layer
    self.inputWire.gradients = d_E_wrt_x
    # print("backward gradient:", d_E_wrt_x)

    self.update()

def update(self):
    global alpha
    # if self.dropout:
    #     self.inputWire.weights -= alpha*(np.outer(self.inputs, self.inputWire.dropout_matrix.dot(self.d_E_wrt_f)))
    # self.inputWire.velocity = self.inputWire.momentum*self.inputWire.velocity - alpha * (np.outer(self.inputs, self.d_E_wrt_f))
    # self.inputWire.weights += self.inputWire.velocity
    self.inputWire.weights -= alpha * (np.outer(self.inputs, self.d_E_wrt_f))

class reluLayer(Layer):
    def __init__(self, inputWire, units, dropout=False):
        super().__init__(inputWire, units, dropout, weight_fn=lambda x, y: np.random.randn(x, y) / np.sqrt(x / 2))

    def forward(self, x, training=False):
        self.inputs = x
        # Add input value of "1" to the input array for the bias weight (w_0*x_0 + ...)
        self.inputs = np.insert(x, 0, 1.0, 0)
        self.linear_outputs = self.inputWire.weights.T.dot(self.inputs)

        self.outputs = self.relu(self.linear_outputs)
        if training and self.dropout:
            self.inputWire.generate_dropout_variables()
            # self.outputs = self.outputs.dot(self.inputWire.dropout_matrix)
            self.outputs = self.outputs * self.inputWire.dropout_vector / self.inputWire.dropout_proba
        return self.outputs

    def relu(self, x):
        # Calculate sigmoid on vector cells
        self.relu_values = np.maximum(0, x)
        return self.relu_values

    def backward(self):
        d_p_wrt_f = np.zeros(self.relu_values.shape)
        d_p_wrt_f[self.relu_values > 0] = 1
        if self.dropout:
            d_p_wrt_f = d_p_wrt_f * self.inputWire.dropout_vector
        self.d_E_wrt_f = np.sum(self.outputWire.gradients[:, 1]) * d_p_wrt_f
        d_E_wrt_x = self.inputWire.weights.dot(
            np.diag(self.d_E_wrt_f)) # Multiply outgoing gradients by the dE/df of their respective node in this layer
        self.inputWire.gradients = d_E_wrt_x
        # print("backward gradient:", d_E_wrt_x)

        self.update()

    def update(self):
        global alpha
        # if self.dropout:
        #     self.inputWire.weights -= alpha*(np.outer(self.inputs, self.inputWire.dropout_matrix.dot(self.d_E_wrt_f)))
        # self.inputWire.velocity = self.inputWire.momentum*self.inputWire.velocity - alpha * (np.outer(self.inputs, self.d_E_wrt_f))
        # self.inputWire.weights += self.inputWire.velocity
        self.inputWire.weights -= alpha * (np.outer(self.inputs, self.d_E_wrt_f))

class softmaxLayer(Layer):
    def __init__(self, inputWire, units, dropout=False):
        super().__init__(inputWire, units, dropout)

    def forward(self, x, training=False):
        self.inputs = x
        # Add input value of "1" to the input array for the bias weight (w_0*x_0 + ...)
        self.inputs = np.insert(x, 0, 1, 0)
        self.outputs = self.softmax(self.inputWire.weights.T.dot(self.inputs))
        return self.outputs

    def softmax(self, x):
        # Calculate sigmoid on vector cells
        exponentials = np.exp(x - np.max(x))
        self.softmax_values = exponentials / np.sum(exponentials)
        return self.softmax_values

    def backward(self):
        self.inputWire.gradients = None

        d_p_wrt_f = self.softmax_values * (1 - self.softmax_values)
        self.d_E_wrt_f = self.outputWire.gradients * d_p_wrt_f # single incoming gradient
        d_E_wrt_x = self.inputWire.weights.dot(

```

```

        np.diag(self.d_E_wrt_f)) # Multiply outgoing gradients by the dE/df of their respective node in this layer
        self.inputWire.gradients = d_E_wrt_x

        self.update()

    def backward(self, y):
        self.inputWire.gradients = None

        self.d_E_wrt_f = self.outputs - y # single incoming gradient
        d_E_wrt_x = self.inputWire.weights.dot(
            np.diag(self.d_E_wrt_f)) # Multiply outgoing gradients by the dE/df of their respective node in this layer
        self.inputWire.gradients = d_E_wrt_x

        self.update()

    def update(self):
        global alpha
        # TODO: Refactor away from the outer() function to allow for mini-batch
        self.inputWire.weights -= alpha * (np.outer(self.inputs, self.d_E_wrt_f))

from numpy.random import multivariate_normal

def unison_shuffled_copies(a, b):
    assert len(a) == len(b)
    p = np.random.permutation(len(a))
    return [a[p], b[p]]

mean_w1 = np.array([0, 0])
covar_w1 = np.array([[1, 0], [0, 1]])
mean_w2 = np.array([1, 0.5])
covar_w2 = np.array([[3, 1], [1, 2]])

# Generate training data and labels
training_w1_data = multivariate_normal(mean_w1, covar_w1, size=(50))
training_w1_labels = np.zeros(training_w1_data.shape[0])
training_w2_data = multivariate_normal(mean_w2, covar_w2, size=(50))
training_w2_labels = np.ones(training_w2_data.shape[0])

# Generate test data and labels
test_w1_data = multivariate_normal(mean_w1, covar_w1, size=(20))
test_w1_labels = np.zeros(test_w1_data.shape[0])
test_w1_labels = test_w1_labels.reshape((-1, 1))
test_w2_data = multivariate_normal(mean_w2, covar_w2, size=(20))
test_w2_labels = np.ones(test_w2_data.shape[0])
test_w2_labels = test_w2_labels.reshape((-1, 1))
training_data = np.r_[training_w1_data, training_w2_data]
training_labels = np.r_[training_w1_labels, training_w2_labels]
training = unison_shuffled_copies(training_data, training_labels)
training[1] = training[1].reshape((-1, 1))
test_data = np.r_[test_w1_data, test_w2_data]
test_labels = np.r_[test_w1_labels, test_w2_labels]
test = unison_shuffled_copies(test_data, test_labels)
test[1] = test[1].reshape((-1, 1))

EPOCHS = 100
nn_architecture = [("input", 2), ("sigmoid", 5), ("sigmoid", 1)]
nn = NeuralNetwork(nn_architecture, dropout=False)
nn.train_and_validate(training, training, EPOCHS)

```