

CSE 575 Homework 1

Andrew Dudley (addudley@asu.edu)

February 5, 2018

Question 1.

Suppose that in your coin flip experiment, you observed a set of α_H heads and α_T tails. Let θ denote the probability of observing heads, whose prior distribution follows $Beta(\beta_H, \beta_T)$, where β_H and β_T are two positive parameters.

- (a) Prove that the posterior distribution $P(\theta|D)$ follows $Beta(\beta_H + \alpha_H, \beta_T + \alpha_T)$

The posterior distribution $P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$. We're told that

$$P(\theta) = \frac{\theta^{\beta_H-1}(1-\theta)^{\beta_T-1}}{B(\beta_H, \beta_T)} \sim Beta(\beta_H, \beta_T),$$

where $B(\beta_H, \beta_T)$ is the beta function.

Note that $P(D|\theta)$ is the likelihood function, which for a Bernoulli experiment is

$$\theta^{\alpha_H}(1-\theta)^{\alpha_T}.$$

Putting these together, we get

$$P(\theta|D) = \frac{\theta^{\beta_H+\alpha_H-1}(1-\theta)^{\beta_T+\alpha_T-1}}{P(D)B(\beta_H, \beta_T)},$$

and because both $P(D)$ and $B(\beta_H, \beta_T)$ are normalizing constants (they don't rely on θ), we can rewrite the equation as

$$P(\theta|D) = \frac{\theta^{\beta_H+\alpha_H-1}(1-\theta)^{\beta_T+\alpha_T-1}}{B(\beta_H + \alpha_H, \beta_T + \alpha_T)},$$

From here, we can see that

$$P(\theta|D) \sim Beta(\beta_H + \alpha_H, \beta_T + \alpha_T).$$

(b) What is the mean of $P(\theta|D)$?

For notational convenience, let $a = \beta_H + \alpha_H$ and $b = \beta_T + \alpha_T$. As they say, statistics is the practice of replacing expectations with averages, so the mean of $P(\theta|D)$ is the same as $\mathbb{E}[P(\theta|D)]$. And, of course, because this value is a probability, we can bound the integral by $[0, 1]$.

$$\mathbb{E}[P(\theta|D)] = \int_0^1 \theta P(\theta|D) d\theta \quad (1)$$

$$= \frac{1}{B(a, b)} \int_0^1 \theta^a (1 - \theta)^{b-1} d\theta \quad (2)$$

Here, we can see that the integral is, itself, a beta function.

$$= \frac{B(a+1, b)}{B(a, b)} \quad (3)$$

$$= \frac{\Gamma(a+1)\Gamma(b)\Gamma(a+b)}{\Gamma(a)\Gamma(b)\Gamma(a+b+1)} \quad (4)$$

$$= \frac{a}{a+b} \quad (5)$$

$$= \frac{\beta_H + \alpha_H}{\beta_H + \alpha_H + \beta_T + \alpha_T} \quad (6)$$

(c) What is the MAP estimator $\hat{\theta}_{MAP}$ of θ ?

$$\begin{aligned} \frac{d}{d\theta} \mathcal{L}(\theta) &= \frac{d}{d\theta} \ln \left(\theta^{\beta_H + \alpha_H - 1} (1 - \theta)^{\beta_T + \alpha_T - 1} \right) \\ &= (\beta_H + \alpha_H - 1) \left[\frac{d}{d\theta} \ln \theta \right] + (\beta_T + \alpha_T - 1) \left[\frac{d}{d\theta} \ln(1 - \theta) \right] \\ &= \frac{\beta_H + \alpha_H - 1}{\theta} - \frac{\beta_T + \alpha_T - 1}{1 - \theta} \end{aligned}$$

And to find a critical point (the minimum value), we set the derivative to 0.

$$\begin{aligned} \frac{\beta_H + \alpha_H - 1}{\theta} - \frac{\beta_T + \alpha_T - 1}{1 - \theta} &= 0 \\ \hat{\theta}_{MAP} &= \frac{\beta_H + \alpha_H - 1}{\beta_H + \alpha_H + \beta_T + \alpha_T - 2} \end{aligned}$$

Question 2.

For this question, assume that $x_1, \dots, x_N \in \mathbb{R}$ are i.i.d. from a normal distribution.

- (a) Let $\hat{\mu}_{MLE}$ denote the MLE of μ . Prove that $\hat{\mu}_{MLE}$ is unbiased.

We'll first need to determine the equation for $\hat{\mu}_{MLE}$, then we'll need to compare its expectation to the population mean μ . Let $D = \{x_1, \dots, x_n\}$. Then, because the data is assumed independent, the likelihood of μ with respect to D can be written as

$$P(D|\mu) = \prod_{i=1}^n p(x_i|\mu)$$

Furthermore, we know that the samples are from a normal distribution, giving us

$$p(x|\mu) = \frac{1}{(2\pi\sigma^2)^{\frac{1}{2}}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

To simplify finding the derivative, take the log of the likelihood function

$$\mathcal{L}(P(D|\mu)) = \sum_{i=1}^n \ln(p(x_i|\mu))$$

Now we'll find the MLE of μ by taking the derivative of the log-likelihood, setting it to 0, and solving for μ

$$\begin{aligned} \frac{d\mathcal{L}}{d\mu} &= \sum_{i=1}^n \frac{d}{d\mu} \ln(p(x_i|\mu)) &&= 0 \\ &= - \sum_{i=1}^n \frac{d}{d\mu} \left(\frac{1}{2} \ln(2\pi\sigma^2) + \frac{1}{2\sigma^2} (x_i - \mu)^2 \right) &&= 0 \\ &= \sum_{i=1}^n (x_i - \hat{\mu}) &&= 0 \\ \hat{\mu} &= \frac{1}{n} \sum_{i=1}^n x_i \end{aligned}$$

Now that we know the equation for $\hat{\mu}$, we'll find its expectation.

$$\begin{aligned}\mathbb{E}[\hat{\mu}] &= \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n x_i\right] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[x_i] \\ &= \frac{1}{n} \sum_{i=1}^n \mu \\ &= \mu\end{aligned}$$

Therefore, $\hat{\mu}_{MLE}$ is unbiased.

- (b) If the true value of μ is unknown, then the MLE estimate of σ^2 is

$$\sigma_{MLE}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu}_{MLE})^2$$

Prove that σ_{MLE}^2 is biased. TODO

Question 3.

- (a) How many independent parameters would there be for the Naive Bayes classifier trained with the given data? What are they?

There would be **thirteen independent parameters**. Note first that RID is a nominal attribute and buys_computer is the dependent (target) attribute, leaving us with age, income, student, and credit_rating as the attributes of interest.

An assumption is made with Naive Bayes that

$$P(X_i | X_1, \dots, i-1, x_{i+1}, \dots, n, Y) = P(X_i | Y) \forall i \in \{1, \dots, n\}$$

A Bayes classifier can be represented as

$$P(Y_c | X_{1..n}) = \frac{P(X_{1..n} | Y_c) P(Y_c)}{\sum_Y P(X_{1..n} | Y) P(Y)}$$

Given the conditional independence assumption of Naive Bayes, this can be simplified to

$$P(Y_c | X_{1..n}) = \frac{\prod_i P(X_i | Y_c) P(Y_c)}{\sum_Y \prod_i P(X_i | Y) P(Y)}$$

Let m_i represent the number of discrete categories in the variable x_i and C represent the number of discrete classes of the target attribute Y . Then, given $P(X = X_{ij}|Y = y_c)$ where $j \in \{1, \dots, m_i\}$ and $c \in \{1, \dots, C\}$, each independent variable x_i contributes $m_i - 1$ independent parameters for each class y_c . Note that, given $m_i - 1$ parameters for the class conditional of an attribute, the final parameter of that class conditional is simply the difference between 1 and the sum of those parameters, and is therefore not independent.

Similarly, we must account for the independent parameters contributed by the prior $P(Y)$, which will be $C - 1$ parameters.

Thus, the equation for the total number of independent parameters in a Naive Bayes model will be

$$C * \sum_{i=1}^n (m_i - 1) + C - 1$$

Plugging in the attributes from the data provided, the number of independent parameters is **thirteen**.

(b)

$$\begin{aligned} P(\text{age} = \text{"youth"} | \text{buys_computer} = \text{"no"}) &= 3/5, P(\text{age} = \text{"middle_aged"} | \text{buys_computer} = \text{"no"}) = 0 \\ P(\text{age} = \text{"youth"} | \text{buys_computer} = \text{"yes"}) &= 2/9, P(\text{age} = \text{"middle_aged"} | \text{buys_computer} = \text{"yes"}) = 4/9 \\ P(\text{income} = \text{"low"} | \text{buys_computer} = \text{"no"}) &= 1/5, P(\text{income} = \text{"medium"} | \text{buys_computer} = \text{"no"}) = 2/5 \\ P(\text{income} = \text{"low"} | \text{buys_computer} = \text{"yes"}) &= 3/9, P(\text{income} = \text{"medium"} | \text{buys_computer} = \text{"yes"}) = 4/9 \\ P(\text{student} = \text{"no"} | \text{buys_computer} = \text{"no"}) &= 4/5 \\ P(\text{student} = \text{"no"} | \text{buys_computer} = \text{"yes"}) &= 3/9 \\ P(\text{credit_rating} = \text{"fair"} | \text{buys_computer} = \text{"no"}) &= 2/5 \\ P(\text{credit_rating} = \text{"fair"} | \text{buys_computer} = \text{"yes"}) &= 6/9 \\ P(\text{buys_computer} = \text{"yes"}) &= 9/14 \end{aligned}$$

(c) Given a new person with features $x = (\text{youth}, \text{medium}, \text{yes}, \text{fair})$, what is $P(Y = \text{yes}|x)$?

$$\frac{\frac{2}{9} * \frac{4}{9} * \frac{6}{9} * \frac{6}{9} * \frac{9}{14}}{\frac{2}{9} * \frac{4}{9} * \frac{6}{9} * \frac{9}{14} + \frac{3}{5} * \frac{2}{5} * \frac{1}{5} * \frac{2}{5} * \frac{5}{14}} = 0.80451$$

Therefore, NB would classify this input at $Y = \text{yes}$.

Question 4.

Suppose we have two positive examples $x_1 = (1, 0)$ and $x_2 = (0, -1)$, and

two negative samples $x_3 = (0, 1)$ and $x_4 = (-1, 0)$. Apply standard gradient ascent method to train a logistic regression classifier (without any regularization term). Initialize the weight vector with two different value and set $w_0^0 = 0$. Would the final weight vector (w^*) be the same for the two different initial values? What are the values? You may assume the learning rate to be a positive real constant η .

For notational and implementation convenience, first append a 1 to each sample vector

$$X = \begin{bmatrix} 1 & 0 & 1 \\ 0 & -1 & 1 \\ 0 & 1 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$